

Problem Statement:

for Insurance dataset implementing all models based on the accuracy, which means which model is highest accuracy. That model is best model for this dataset. to finding best model

In [2]:

```
import numpy as np, pandas as pd, matplotlib.pyplot as plt, seaborn as sns
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

Data collection:

In [3]:

```
df=pd.read_csv(r"C:\Users\raja\Downloads\insurance (1).csv")
df
```

Out[3]:

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|--------|--------|----------|--------|-----------|-------------|
| 0 | 19 | female | 27.900 | 0 | yes | southwest | 16884.92400 |
| 1 | 18 | male | 33.770 | 1 | no | southeast | 1725.55230 |
| 2 | 28 | male | 33.000 | 3 | no | southeast | 4449.46200 |
| 3 | 33 | male | 22.705 | 0 | no | northwest | 21984.47061 |
| 4 | 32 | male | 28.880 | 0 | no | northwest | 3866.85520 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1333 | 50 | male | 30.970 | 3 | no | northwest | 10600.54830 |
| 1334 | 18 | female | 31.920 | 0 | no | northeast | 2205.98080 |
| 1335 | 18 | female | 36.850 | 0 | no | southeast | 1629.83350 |
| 1336 | 21 | female | 25.800 | 0 | no | southwest | 2007.94500 |
| 1337 | 61 | female | 29.070 | 0 | yes | northwest | 29141.36030 |

1338 rows × 7 columns

Data cleaning and Preprocessing

In [4]:

```
df.head()
```

Out[4]:

| | age | sex | bmi | children | smoker | region | charges |
|---|-----|--------|--------|----------|--------|-----------|-------------|
| 0 | 19 | female | 27.900 | 0 | yes | southwest | 16884.92400 |
| 1 | 18 | male | 33.770 | 1 | no | southeast | 1725.55230 |
| 2 | 28 | male | 33.000 | 3 | no | southeast | 4449.46200 |
| 3 | 33 | male | 22.705 | 0 | no | northwest | 21984.47061 |
| 4 | 32 | male | 28.880 | 0 | no | northwest | 3866.85520 |

In [5]:

```
df.tail()
```

Out[5]:

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|--------|-------|----------|--------|-----------|------------|
| 1333 | 50 | male | 30.97 | 3 | no | northwest | 10600.5483 |
| 1334 | 18 | female | 31.92 | 0 | no | northeast | 2205.9808 |
| 1335 | 18 | female | 36.85 | 0 | no | southeast | 1629.8335 |
| 1336 | 21 | female | 25.80 | 0 | no | southwest | 2007.9450 |
| 1337 | 61 | female | 29.07 | 0 | yes | northwest | 29141.3603 |

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [7]:

```
df.describe()
```

Out[7]:

| | age | bmi | children | charges |
|-------|-------------|-------------|-------------|--------------|
| count | 1338.000000 | 1338.000000 | 1338.000000 | 1338.000000 |
| mean | 39.207025 | 30.663397 | 1.094918 | 13270.422265 |
| std | 14.049960 | 6.098187 | 1.205493 | 12110.011237 |
| min | 18.000000 | 15.960000 | 0.000000 | 1121.873900 |
| 25% | 27.000000 | 26.296250 | 0.000000 | 4740.287150 |
| 50% | 39.000000 | 30.400000 | 1.000000 | 9382.033000 |
| 75% | 51.000000 | 34.693750 | 2.000000 | 16639.912515 |
| max | 64.000000 | 53.130000 | 5.000000 | 63770.428010 |

In [8]:

```
df.size
```

Out[8]:

9366

In [9]:

```
df.shape
```

Out[9]:

(1338, 7)

In [10]:

```
df.isna().any()
```

Out[10]:

```
age      False
sex      False
bmi      False
children False
smoker   False
region   False
charges  False
dtype: bool
```

In [11]:

```
df['sex'].value_counts()
```

Out[11]:

```
sex
male      676
female    662
Name: count, dtype: int64
```

In [12]:

```
df['smoker'].value_counts()
```

Out[12]:

```
smoker
no      1064
yes      274
Name: count, dtype: int64
```

In [13]:

```
convert={"region":{"southeast":1,"southwest":2,"northwest":3,"northeast":4}}
df=df.replace(convert)
df
```

Out[13]:

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|--------|--------|----------|--------|--------|-------------|
| 0 | 19 | female | 27.900 | 0 | yes | 2 | 16884.92400 |
| 1 | 18 | male | 33.770 | 1 | no | 1 | 1725.55230 |
| 2 | 28 | male | 33.000 | 3 | no | 1 | 4449.46200 |
| 3 | 33 | male | 22.705 | 0 | no | 3 | 21984.47061 |
| 4 | 32 | male | 28.880 | 0 | no | 3 | 3866.85520 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1333 | 50 | male | 30.970 | 3 | no | 3 | 10600.54830 |
| 1334 | 18 | female | 31.920 | 0 | no | 4 | 2205.98080 |
| 1335 | 18 | female | 36.850 | 0 | no | 1 | 1629.83350 |
| 1336 | 21 | female | 25.800 | 0 | no | 2 | 2007.94500 |
| 1337 | 61 | female | 29.070 | 0 | yes | 3 | 29141.36030 |

1338 rows × 7 columns

In [14]:

```
df['region'].value_counts()
```

Out[14]:

```
region
1      364
2      325
3      325
4      324
Name: count, dtype: int64
```

In [15]:

```
convert={"sex":{"male":0,"female":1}}
df=df.replace(convert)
df
```

Out[15]:

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|-----|--------|----------|--------|--------|-------------|
| 0 | 19 | 1 | 27.900 | 0 | yes | 2 | 16884.92400 |
| 1 | 18 | 0 | 33.770 | 1 | no | 1 | 1725.55230 |
| 2 | 28 | 0 | 33.000 | 3 | no | 1 | 4449.46200 |
| 3 | 33 | 0 | 22.705 | 0 | no | 3 | 21984.47061 |
| 4 | 32 | 0 | 28.880 | 0 | no | 3 | 3866.85520 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1333 | 50 | 0 | 30.970 | 3 | no | 3 | 10600.54830 |
| 1334 | 18 | 1 | 31.920 | 0 | no | 4 | 2205.98080 |
| 1335 | 18 | 1 | 36.850 | 0 | no | 1 | 1629.83350 |
| 1336 | 21 | 1 | 25.800 | 0 | no | 2 | 2007.94500 |
| 1337 | 61 | 1 | 29.070 | 0 | yes | 3 | 29141.36030 |

1338 rows × 7 columns

In [16]:

```
convert={"smoker":{"no":1,"yes":2}}  
df=df.replace(convert)  
df
```

Out[16]:

| | age | sex | bmi | children | smoker | region | charges |
|------|-----|-----|--------|----------|--------|--------|-------------|
| 0 | 19 | 1 | 27.900 | 0 | 2 | 2 | 16884.92400 |
| 1 | 18 | 0 | 33.770 | 1 | 1 | 1 | 1725.55230 |
| 2 | 28 | 0 | 33.000 | 3 | 1 | 1 | 4449.46200 |
| 3 | 33 | 0 | 22.705 | 0 | 1 | 3 | 21984.47061 |
| 4 | 32 | 0 | 28.880 | 0 | 1 | 3 | 3866.85520 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1333 | 50 | 0 | 30.970 | 3 | 1 | 3 | 10600.54830 |
| 1334 | 18 | 1 | 31.920 | 0 | 1 | 4 | 2205.98080 |
| 1335 | 18 | 1 | 36.850 | 0 | 1 | 1 | 1629.83350 |
| 1336 | 21 | 1 | 25.800 | 0 | 1 | 2 | 2007.94500 |
| 1337 | 61 | 1 | 29.070 | 0 | 2 | 3 | 29141.36030 |

1338 rows × 7 columns

In [17]:

```
features=df.columns[0:5]  
target=df.columns[-1]
```

Exploratory Data Analysis


Data visualization

In [19]:

```
sns.displot(df["charges"])
```

Out[19]:

<seaborn.axisgrid.FacetGrid at 0x1af445eec50>



In [20]:

```
features=df.columns[0:6]  
features
```

Out[20]:

Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region'], dtype='object')

In [21]:

```
target=df.columns[-1]  
target
```

Out[21]:

'charges'

In [22]:

```
#training our model
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
x=df[features].values
y=df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.50,random_state=17)
x_train.shape
x_test.shape
scaler=StandardScaler()
x_train=scaler.fit_transform(x_train)
x_test=scaler.fit_transform(x_test)
```

Data Modelling

Linear Regression

In [23]:

```
from sklearn.linear_model import LinearRegression
reg=LinearRegression()
```

In [24]:

```
reg.fit(x_train,y_train)
```

Out[24]:

LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.


In [25]:

```
print(reg.score(x_test,y_test))
```

0.7553050021744551

In [26]:

```
y_pred=reg.predict(x_test)
plt.scatter(y_test,y_pred)
plt.show()
```



In [27]:

```
#evaluation of model
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

In [28]:

```
model=LinearRegression()
model.fit(x_train,y_train)
```

Out[28]:

LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [29]:

```
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.7553050021744551

RIDGE

In [30]:

```
from sklearn.linear_model import Ridge,Lasso,ElasticNet
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
```

In [31]:

```
#ridge regression model
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
print("\nRidge Model\n")
print("Train Score for ridge model is",(train_score_ridge))
print("Test Score for ridge model is",(test_score_ridge))
```


Ridge Model

Train Score for ridge model is 0.7408189083612313

Test Score for ridge model is 0.7555036308713557

In [32]:

```
plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color="blue")
plt.plot(features,reg.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color="green")
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



LASSO Regression

In [33]:

```
#Elastic Net
re=ElasticNet()
re.fit(x,y)
print(re.coef_)
#print(re.intercept_)
```

```
[ 244.53200843 -323.29321608  327.88975459  389.15389489 5841.24528723
  75.2866119 ]
```

In [34]:

```
#Lasso regression
lassoReg=Lasso(alpha=10)
lassoReg.fit(x_train,y_train)
train_score_lasso=lassoReg.score(x_train,y_train)
test_score_lasso=lassoReg.score(x_test,y_test)
print("\nLasso Model\n")
print("Train Score for lasso model is",(train_score_lasso))
print("Test Score for lasso model is",(test_score_lasso))
```

lasso Model

Train Score for lasso model is 0.7409768055935853

Test Score for lasso model is 0.7553619179215416

ElasticNet

In [35]:

```
re.score(x,y)
```

Out[35]:

0.3919002752780354

In [36]:

```
y_pred_elastic=re.predict(x_train)
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 830429140.4997891

Logistic Regression

In [37]:

```
from sklearn.linear_model import LogisticRegression
```

In [38]:

```
pd.set_option("display.max_rows",10000000000)
pd.set_option('display.max_columns',10000000000)
pd.set_option('display.width',95)
```

In [39]:

```
print("this dataframe has %d rows and %d columns"%(df.shape))
```

this dataframe has 1338 rows and 7 columns

In [40]:

```
df.head()
```

Out[40]:

| | age | sex | bmi | children | smoker | region | charges |
|---|-----|-----|--------|----------|--------|--------|-------------|
| 0 | 19 | 1 | 27.900 | 0 | 2 | 2 | 16884.92400 |
| 1 | 18 | 0 | 33.770 | 1 | 1 | 1 | 1725.55230 |
| 2 | 28 | 0 | 33.000 | 3 | 1 | 1 | 4449.46200 |
| 3 | 33 | 0 | 22.705 | 0 | 1 | 3 | 21984.47061 |
| 4 | 32 | 0 | 28.880 | 0 | 1 | 3 | 3866.85520 |

In [41]:

```
features_matrix=df[["age","sex","bmi","region"]]  
features_matrix.columns=["age","sex","bmi","region"]  
target_matrix=df.iloc[:,-3]
```

In [42]:

```
print('The features matrix has %d Rows and %d columns'%(features_matrix.shape))  
print('The features matrix has %d Rows and %d columns'%(np.array(target_matrix).reshape(
```

The features matrix has 1338 Rows and 4 columns

The features matrix has 1338 Rows and 1 columns

In [43]:

```
features_matrix_Standardized=StandardScaler().fit_transform(features_matrix)
```

In [44]:

```
algorithm=LogisticRegression(max_iter=100)
```

In [45]:

```
Logistic_Regression_model=algorithm.fit(features_matrix_Standardized,target_matrix)
```

In [46]:

```
observation=[[28,1,28.880,2]]
```

In [47]:

```
predictions=Logistic_Regression_model.predict(observation)  
print('The model predicted the observtaion to belong to class %s'%(predictions))  
print('The algorithm was trained to predict one of the two calsses : %s'%(algorithm.clas
```

The model predicted the observtaion to belong to class [1]

The algorithm was trained to predict one of the two calsses : [1 2]

In [48]:

```
print("""The model says the prbability of the observation we passed belonging to calss [
print()
print("""The model says the prbability of the observation we passed belonging to calss [
```

The model says the prbability of the observation we passed belonging to ca
lss ['0'] Is 0.9800486899442347

The model says the prbability of the observation we passed belonging to ca
lss ['1'] Is 0.019951310055765316

In [49]:

```
x1=np.array(df['charges']).reshape(-1,1)
```

In [50]:


```
sns.regplot(x=x1,y=target_matrix,data=df,logistic=True,ci=None)
```

C:\Users\raja\AppData\Local\Programs\Python\Python310\lib\site-packages\st
atsmodels\genmod\amilies\links.py:198: RuntimeWarning: overflow encounter
ed in exp

```
t = np.exp(-z)
```

Out[50]:

<Axes: ylabel='smoker'>



Decision trees

In [51]:

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
```

In [52]:

df

Out[52]:

| | age | sex | bmi | children | smoker | region | charges |
|----|-----|-----|--------|----------|--------|--------|--------------|
| 0 | 19 | 1 | 27.900 | 0 | 2 | 2 | 16884.924000 |
| 1 | 18 | 0 | 33.770 | 1 | 1 | 1 | 1725.552300 |
| 2 | 28 | 0 | 33.000 | 3 | 1 | 1 | 4449.462000 |
| 3 | 33 | 0 | 22.705 | 0 | 1 | 3 | 21984.470610 |
| 4 | 32 | 0 | 28.880 | 0 | 1 | 3 | 3866.855200 |
| 5 | 31 | 1 | 25.740 | 0 | 1 | 1 | 3756.621600 |
| 6 | 46 | 1 | 33.440 | 1 | 1 | 1 | 8240.589600 |
| 7 | 37 | 1 | 27.740 | 3 | 1 | 3 | 7281.505600 |
| 8 | 37 | 0 | 29.830 | 2 | 1 | 4 | 6406.410700 |
| 9 | 60 | 1 | 25.840 | 0 | 1 | 3 | 28923.136920 |
| 10 | 25 | 0 | 26.220 | 0 | 1 | 4 | 2721.320800 |

In [53]:

```
df["sex"].value_counts()
```

Out[53]:

```
sex
0    676
1    662
Name: count, dtype: int64
```

In [54]:

```
df["smoker"].value_counts()
```

Out[54]:

```
smoker
1    1064
2     274
Name: count, dtype: int64
```


In [55]:

```
x=["age","sex","children","bmi"]
y=["0","1"]
all_inputs=df[x]
all_classes=df["smoker"]
```

In [56]:

```
x_train,x_test,y_train,y_test=train_test_split(all_inputs,all_classes,test_size=0.25)
x_train.shape,x_test.shape
```

Out[56]:

```
((1003, 4), (335, 4))
```

In [57]:

```
clf=DecisionTreeClassifier(random_state=0)
```

In [58]:

```
clf.fit(x_train,y_train)
```

Out[58]:

```
DecisionTreeClassifier(random_state=0)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [59]:

```
score=clf.score(x_test,y_test)
```

In [60]:

```
print(score)
```

```
0.6895522388059702
```

In [61]:

```
clf.score(x_train,y_train)
```

Out[61]:

```
0.9990029910269193
```

Random Forest

In [62]:

```
df1=df.drop("charges",axis=1)
df1
```

Out[62]:

| | age | sex | bmi | children | smoker | region |
|----|-----|-----|--------|----------|--------|--------|
| 0 | 19 | 1 | 27.900 | 0 | 2 | 2 |
| 1 | 18 | 0 | 33.770 | 1 | 1 | 1 |
| 2 | 28 | 0 | 33.000 | 3 | 1 | 1 |
| 3 | 33 | 0 | 22.705 | 0 | 1 | 3 |
| 4 | 32 | 0 | 28.880 | 0 | 1 | 3 |
| 5 | 31 | 1 | 25.740 | 0 | 1 | 1 |
| 6 | 46 | 1 | 33.440 | 1 | 1 | 1 |
| 7 | 37 | 1 | 27.740 | 3 | 1 | 3 |
| 8 | 37 | 0 | 29.830 | 2 | 1 | 4 |
| 9 | 60 | 1 | 25.840 | 0 | 1 | 3 |
| 10 | 25 | 0 | 26.220 | 0 | 1 | 4 |

In [63]:

```
df2=df1.drop("children",axis=1)
df2
```

Out[63]:

| | age | sex | bmi | smoker | region |
|----|-----|-----|--------|--------|--------|
| 0 | 19 | 1 | 27.900 | 2 | 2 |
| 1 | 18 | 0 | 33.770 | 1 | 1 |
| 2 | 28 | 0 | 33.000 | 1 | 1 |
| 3 | 33 | 0 | 22.705 | 1 | 3 |
| 4 | 32 | 0 | 28.880 | 1 | 3 |
| 5 | 31 | 1 | 25.740 | 1 | 1 |
| 6 | 46 | 1 | 33.440 | 1 | 1 |
| 7 | 37 | 1 | 27.740 | 1 | 3 |
| 8 | 37 | 0 | 29.830 | 1 | 4 |
| 9 | 60 | 1 | 25.840 | 1 | 3 |
| 10 | 25 | 0 | 26.220 | 1 | 4 |

In [64]:

```
x=df2.drop("smoker",axis=1)
y=df2["smoker"]
```

In [65]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=0)
x_train.shape,x_test.shape
```

Out[65]:

```
((1003, 4), (335, 4))
```

In [66]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[66]:

```
RandomForestClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [67]:

```
rf=RandomForestClassifier()
```

In [68]:

```
params={"max_depth":[2,3,5,10,20], 'min_samples_leaf':[5,10,20,50,100,200], 'n_estimators'
```

In [69]:

```
from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[69]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [2, 3, 5, 10, 20],
                          'min_samples_leaf': [5, 10, 20, 50, 100, 200],
                          'n_estimators': [10, 25, 30, 50, 100, 200]},
             scoring='accuracy')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [70]:

```
grid_search.best_score_
```

Out[70]:

0.7996039792924112

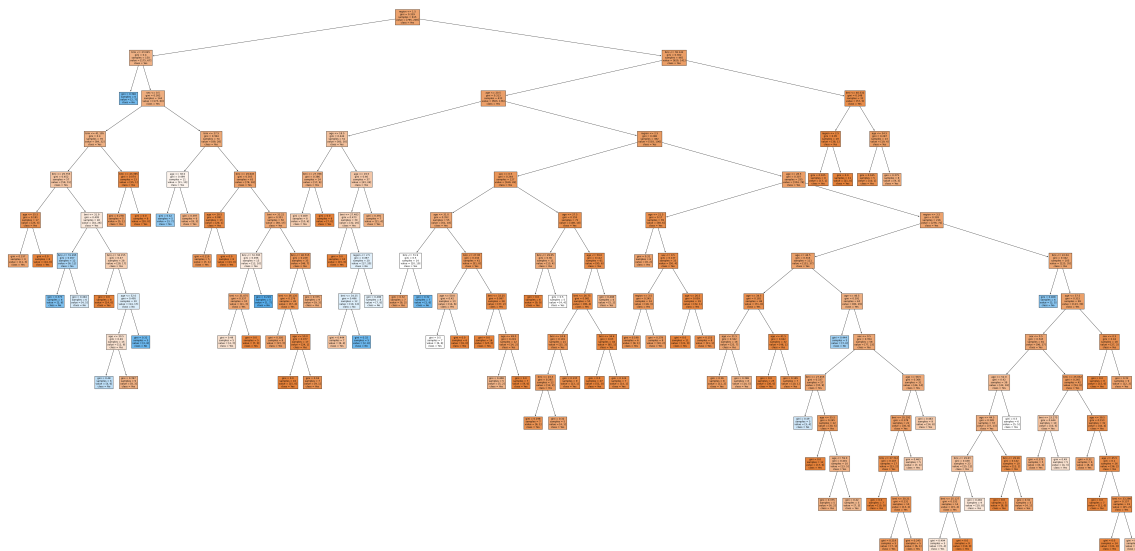
In [71]:

```
rf_best=grid_search.best_estimator_  
print(rf_best)
```

RandomForestClassifier(max_depth=20, min_samples_leaf=5, n_estimators=25)

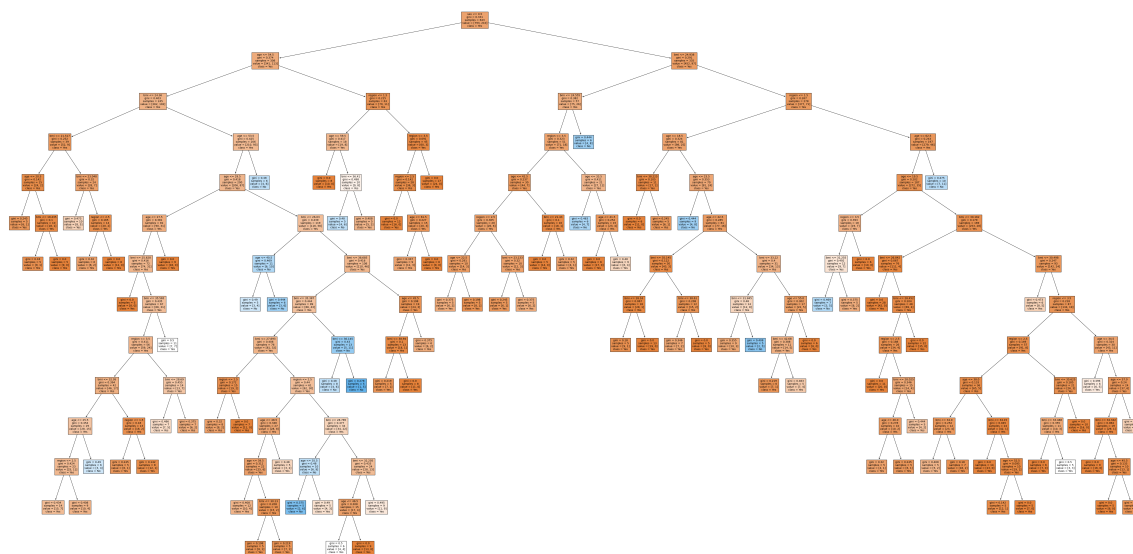
In [72]:

```
from sklearn.tree import plot_tree  
plt.figure(figsize=(80,40))  
plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],filled
```



In [73]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=['Yes','No'],filled
```



In [74]:

```
rf_best.feature_importances_
```

Out[74]:

```
array([0.39397304, 0.0469756 , 0.46710501, 0.09194635])
```

In [75]:

```
imp_df=pd.DataFrame({"Varname":x_train.columns,"Imp":rf_best.feature_importances_})
imp_df.sort_values(by="Imp",ascending=False)
```

Out[75]:

| | Varname | Imp |
|---|---------|----------|
| 2 | bmi | 0.467105 |
| 0 | age | 0.393973 |
| 3 | region | 0.091946 |
| 1 | sex | 0.046976 |

In [83]:

```
print(rfc.score(x_train,y_train))
```

```
0.9990029910269193
```

In [84]:

```
print(rfc.score(x_test,y_test))
```

```
0.7044776119402985
```

Conclusion:

After implementing all models in the dataset RandomForest got highest accuracy. so in this dataset RandomForest is the best model